

# Data Science and Business Intel Project

```
library(tidyverse)
library(ggplot2)
library(ggROC)
library(cowplot)
library(reshape2)
library(car)
library(caret)
library(leaps)
library(bestglm)
library(plotly)
library(webshot)
library(DataExplorer)
library(purrr)
library(rpart)
library(rpart.plot)
library(randomForest)
library(e1071)
library(pROC)
source("../[5]Script/Confusion_matrix.R")
source("../[5]Script/cutoff.R")

# Read in raw data
ds <- (read.csv("../[4]source/WA_Fn-UseC_-Telco-Customer-Churn.csv"))
ds$SeniorCitizen <- as.factor(ds$SeniorCitizen)
#####

# Metadata
##### Customer
# Service that each customer has signed up for
# Demographic information
# Customer Account Information
#####
# Data Type:
#####
# 16 Categorical Variables:
# - 6 Binary Variables (Gender, Senior Citizen, Partner, Dependents, Phone Service, Paperless L
# - 9 3-Factor level Variable (Multiple Lines, Internet Service, Online Security, Online Backu
# - 1 4-Factor level Variable (Payment Method)
#####
# 3 Continious Variables:
# - Tenure, Monthly Charge, Total Charge
#####
# 1 Target Variables:
# - Churn
#####
```

## Data Cleaning

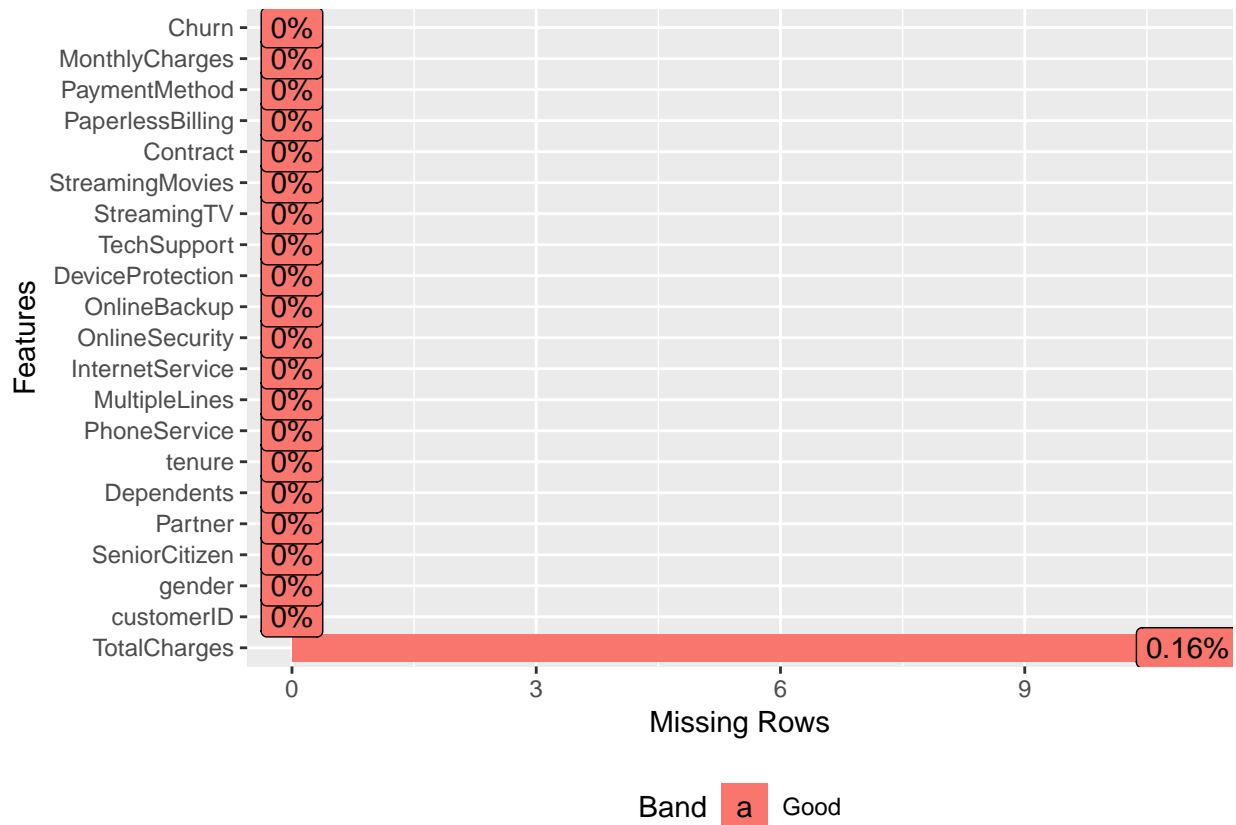
```
# missing data analysis
# 1) We first check if missing exist within our dataset
print(paste0("The dataset contains missing data: ", any(is.na(ds))))

## [1] "The dataset contains missing data: TRUE"

if (any(is.na(ds)) == "TRUE"){
  print(paste0("The total number of missing data(s) are: ", sum(is.na(ds))))
  print(paste0("The variable(s) with missing data(s) are: ", colnames(ds)[colSums(is.na(ds))])
}

## [1] "The total number of missing data(s) are: 11"
## [1] "The variable(s) with missing data(s) are: TotalCharges"

plot_missing(ds)
```



```
# 2) Filter the missing data into a its own dataset for further analysis
df_na <- ds[rowSums(is.na(ds))>0,]
df_na[c("gender", "tenure", "PhoneService", "InternetService", "Contract", "MonthlyCharges", "TotalCharges")]

##      gender tenure PhoneService InternetService Contract MonthlyCharges
## 489  Female      0           No             DSL Two year      52.55
## 754   Male      0           Yes             No Two year      20.25
```

##	937	Female	0	Yes	DSL Two year	80.85
##	1083	Male	0	Yes	No Two year	25.75
##	1341	Female	0	No	DSL Two year	56.05
##	3332	Male	0	Yes	No Two year	19.85
##	3827	Male	0	Yes	No Two year	25.35
##	4381	Female	0	Yes	No Two year	20.00
##	5219	Male	0	Yes	No One year	19.70
##	6671	Female	0	Yes	DSL Two year	73.35
##	6755	Male	0	Yes	DSL Two year	61.90
##		TotalCharges	Churn			
##	489	NA	No			
##	754	NA	No			
##	937	NA	No			
##	1083	NA	No			
##	1341	NA	No			
##	3332	NA	No			
##	3827	NA	No			
##	4381	NA	No			
##	5219	NA	No			
##	6671	NA	No			
##	6755	NA	No			

From the above missing data analysis, we are able to see out of the 7043 observation of 21 variables there are only 11 missing values and they are belong to the TOTAL CHARGES column(.16%), hence we are working with a pretty clean dataset.

An possible explanation for this mssing values is: (1) These customer never paid anything to the company (2) Tenure for all these customer are 0, thus meaning that this may be their first month with the company and thus the company hasn't charged them.

For these 11 missing data, we can either: (1) Impute the total charge value (2) Set total charge value to be zero (3) Remove them from the data set

Since we have a relatively large dataset, and that none of the customer with missing value have churn, thus for convience of the analysis, we will drop the 11 observation with missing TOTAL CHARGE. ## Data Exploration

```
df_clean <- ds %>%
  na.omit() %>%
  select(-1)

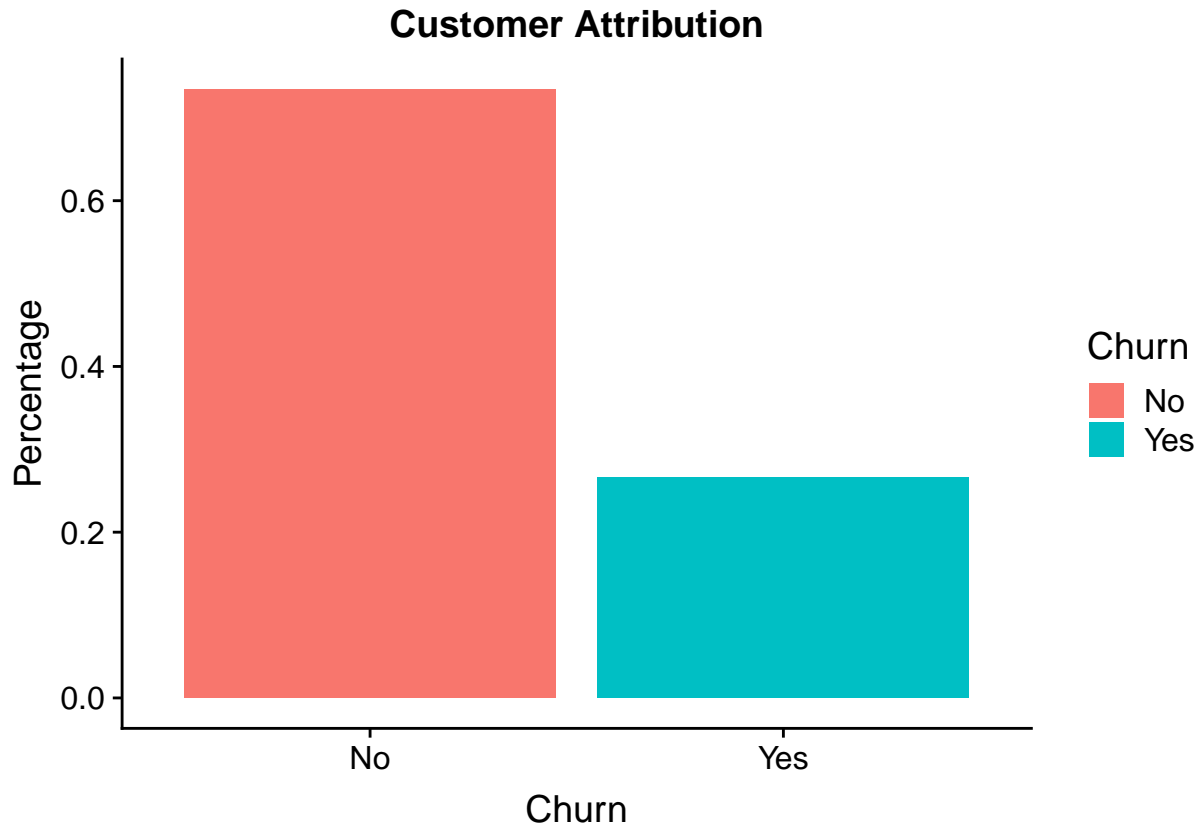
summary(df_clean)
```

##	gender	SeniorCitizen	Partner	Dependents	tenure
##	Female:3483	0:5890	No :3639	No :4933	Min. : 1.00
##	Male :3549	1:1142	Yes:3393	Yes:2099	1st Qu.: 9.00
##					Median :29.00
##					Mean :32.42
##					3rd Qu.:55.00
##					Max. :72.00

```
## PhoneService      MultipleLines      InternetService
## No : 680      No      :3385      DSL      :2416
## Yes:6352      No phone service: 680      Fiber optic:3096
##      Yes      :2967      No      :1520
##
##
##
##      OnlineSecurity      OnlineBackup
## No      :3497      No      :3087
## No internet service:1520      No internet service:1520
## Yes      :2015      Yes      :2425
##
##
##
##      DeviceProtection      TechSupport
## No      :3094      No      :3472
## No internet service:1520      No internet service:1520
## Yes      :2418      Yes      :2040
##
##
##
##      StreamingTV      StreamingMovies
## No      :2809      No      :2781
## No internet service:1520      No internet service:1520
## Yes      :2703      Yes      :2731
##
##
##
##      Contract      PaperlessBilling      PaymentMethod
## Month-to-month:3875      No :2864      Bank transfer (automatic):1542
## One year      :1472      Yes:4168      Credit card (automatic) :1521
## Two year      :1685      Electronic check      :2365
##      Mailed check      :1604
##
##
## MonthlyCharges      TotalCharges      Churn
## Min.      : 18.25      Min.      : 18.8      No :5163
## 1st Qu.: 35.59      1st Qu.: 401.4      Yes:1869
## Median : 70.35      Median :1397.5
## Mean      : 64.80      Mean      :2283.3
## 3rd Qu.: 89.86      3rd Qu.:3794.7
## Max.      :118.75      Max.      :8684.8
```

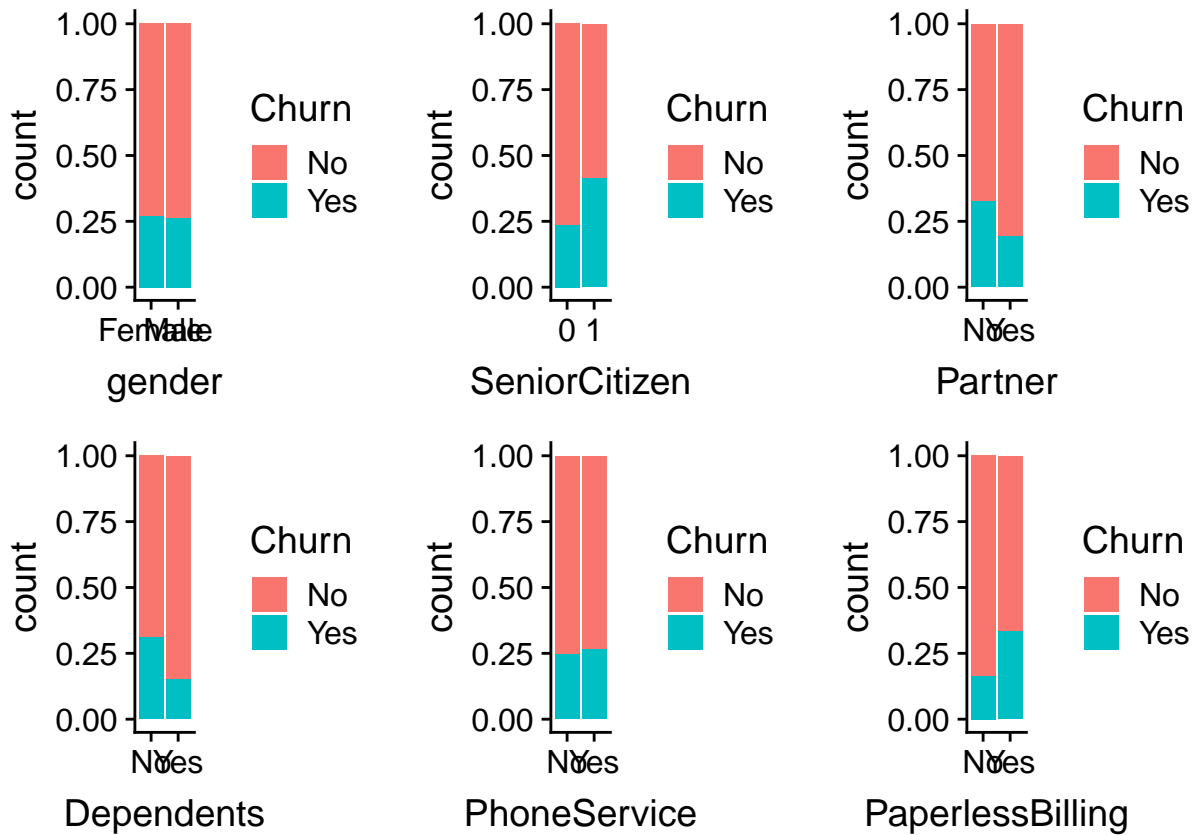
```
# Binary variable distribution in Customer attribution
```

```
ggplot(data = df_clean, aes(x = Churn, y = (..count..)/sum(..count..), fill = Churn))+
  geom_bar()+
  ggtitle("Customer Attribution")+
  ylab("Percentage")
```

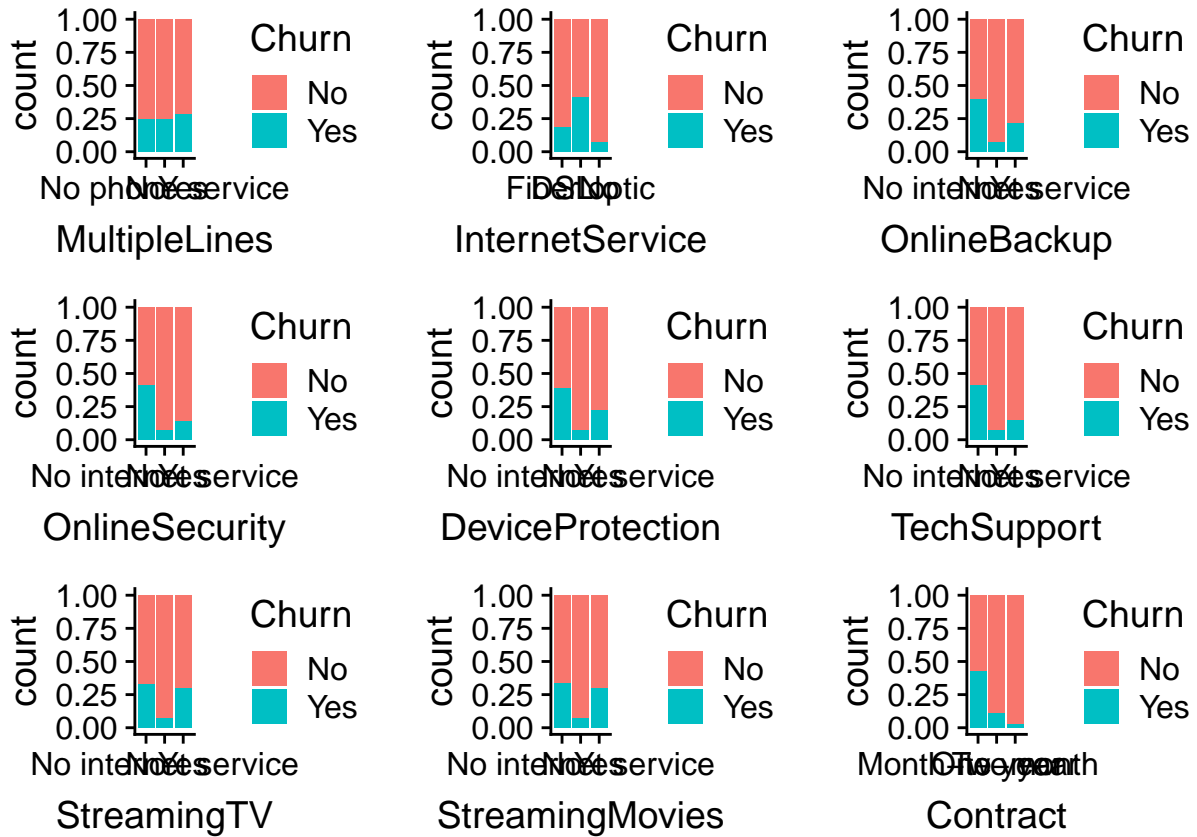


Of our dataset, 26% of the customer has left the platform within the past month

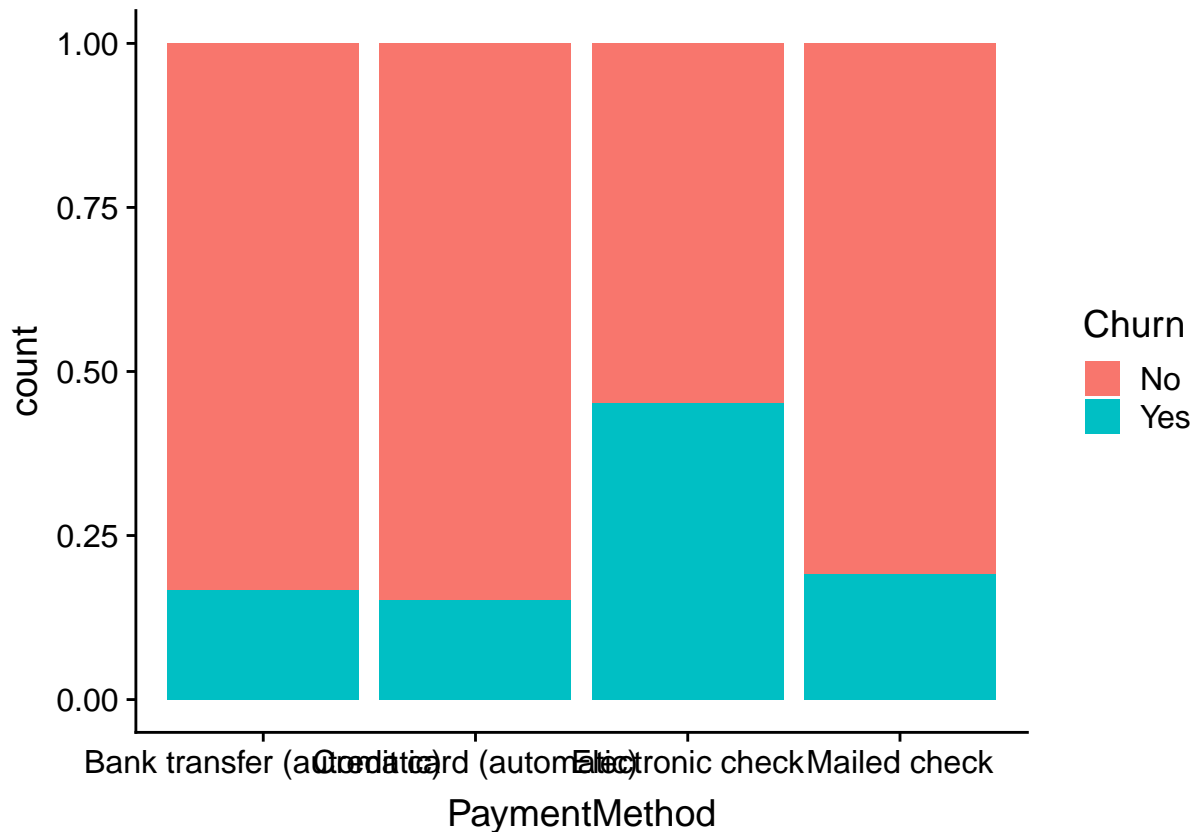
```
# Categorical Variable Analysis  
# Binary binary variables Analysis  
options(repr.plot.width = 12, repr.plot.height = 8)  
plot_grid(  
  ggplot(data = df_clean, aes(gender, fill = Churn))+geom_bar(position = "fill"),  
  ggplot(data = df_clean, aes(SeniorCitizen, fill = Churn))+geom_bar(position = "fill"),  
  ggplot(data = df_clean, aes(Partner, fill = Churn))+geom_bar(position = "fill"),  
  ggplot(data = df_clean, aes(Dependents, fill = Churn))+geom_bar(position = "fill"),  
  ggplot(data = df_clean, aes(PhoneService, fill = Churn))+geom_bar(position = "fill"),  
  ggplot(data = df_clean, aes(PaperlessBilling, fill = Churn))+geom_bar(position = "fill")  
)
```



```
plot_grid(
  ggplot(data = df_clean, aes(MultipleLines, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(InternetService, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(OnlineBackup, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(OnlineSecurity, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(DeviceProtection, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(TechSupport, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(StreamingTV, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(StreamingMovies, fill = Churn))+geom_bar(position = "fill"),
  ggplot(data = df_clean, aes(Contract, fill = Churn))+geom_bar(position = "fill")
)
```



```
ggplot(data = df_clean, aes(x=PaymentMethod, fill=Churn))+
  geom_bar(position = "fill")
```

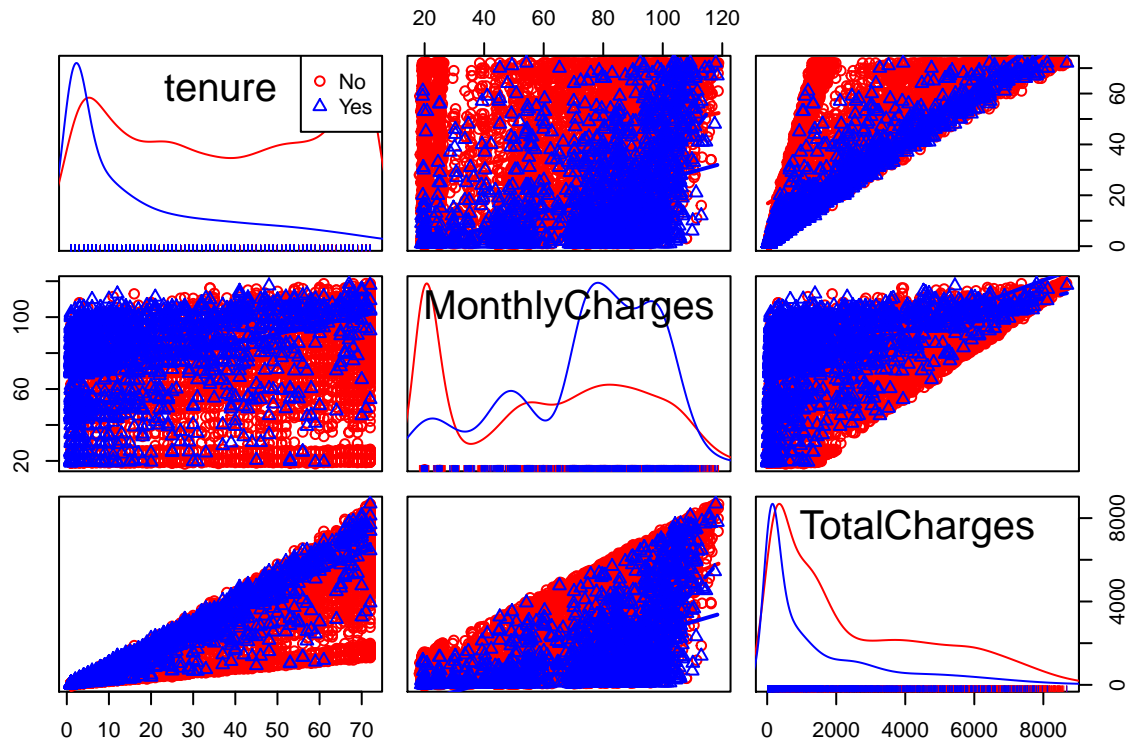


Some Trends - Senior Citizens churn percentage are higher - Customer with dependents or partners tend to have lower churn rate compared to counterparts - Customer with paperless billing have higher churn rate - Customer with Fiber Optic Internet Service have significant higher churn rate - Customer with No online security, or online backup or tech support have higher churn rate - Customer with monthly subscription are more likely to churn compared to customer with one- or two-year contract - Customer with Electronic Check payment method tend to leave our client more compared to other options.

```
## Continous Variable Analysis
# p <- plot_ly(df_clean,
#             x = ~MonthlyCharges,
#             y = ~TotalCharges,
#             z = ~tenure,
#             color = ~Churn,
#             marker = list(
#               size = 2)) %>%
#             add_markers() %>%
#             layout(scene = list(
#               xaxis = list(title = "Monthly Charges"),
#               yaxis = list(title = "Total Charges"),
#               zaxis = list(title = "Tenure")
#             ))
# p
```



```
# Correlation matrix of continuous variable analysis (thank you very much)
scatterplotMatrix(~ tenure + MonthlyCharges + TotalCharges|Churn, data = df_clean, col = c("red", "blue"))
```




Placeholder for observation for the continuous

```
# Data Cleaning and Standardization
df_clean2 <- df_clean %>%
  mutate(MultipleLines=replace(MultipleLines,MultipleLines=="No phone service", "No"))
  mutate(OnlineSecurity=replace(OnlineSecurity,OnlineSecurity=="No internet service", "No"))
  mutate(DeviceProtection=replace(DeviceProtection,DeviceProtection=="No internet service", "No"))
  mutate(TechSupport=replace(TechSupport,TechSupport=="No internet service","No"))
  mutate(StreamingTV=replace(StreamingTV,StreamingTV=="No internet service","No"))
  mutate(StreamingMovies=replace(StreamingMovies,StreamingMovies=="No internet service","No"))
  mutate(OnlineBackup=replace(OnlineBackup,OnlineBackup=="No internet service","No"))
  mutate(tenure=scale(tenure)) %>%
  mutate(MonthlyCharges=scale(MonthlyCharges)) %>%
  mutate(TotalCharges=scale(TotalCharges))

# Feature Importance Analysis
# generateFilterValuesData(task, "randomForest.importance")
df_clean2.feature <- randomForest(Churn~., data = df_clean2, importance = FALSE, ntree = 500, mtry = 5)

varImpPlot(df_clean2.feature)
```

## df\_clean2.feature



```
# Split data into training and validation split
set.seed(1994)
training <- sample(2,nrow(df_clean2),replace=TRUE,prob=c(.8,.2))

# GLM Analysis
# Still need to look at threshold analysis
# Full GLM
df_clean.fulllogit <- glm(Churn~.,
                          family = binomial,
                          data = df_clean2[training==1,])

getinfo(df_clean.fulllogit,df_clean2)[c("confusion_matrix", "accuracy","sensitivity")]

## $confusion_matrix
##      predicted
## observed  0   1
##      No  770 251
##      Yes  84 278
##
## $accuracy
## [1] 0.757773
##
## $sensitivity
## [1] 0.7679558
```

```

# Using Forward Approach to search for GLM model with lowest BIC

# tmp.modelsearch <- bestglm(df_clean2[training==1,],IC = "BIC", family = binomial, method = "BIC")
# tmp.modelsearch$BestModels
# tmp.modelsearch$BestModel

# Best GLM Model
df_clean.bestlogit <- glm(Churn~
                          SeniorCitizen +
                          tenure +
                          PhoneService +
                          InternetService +
                          OnlineSecurity +
                          Contract +
                          PaperlessBilling +
                          PaymentMethod +
                          TotalCharges,
                          family = binomial,
                          data = df_clean2[training==1,])

summary(df_clean.bestlogit)

```

```

##
## Call:
## glm(formula = Churn ~ SeniorCitizen + tenure + PhoneService +
##      InternetService + OnlineSecurity + Contract + PaperlessBilling +
##      PaymentMethod + TotalCharges, family = binomial, data = df_clean2[training ==
##      1, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7854  -0.6901  -0.2898   0.7639   3.5287
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.15482    0.16519  -6.991 2.74e-12
## SeniorCitizen1     0.36347    0.09087   4.000 6.33e-05
## tenure          -1.60719    0.16355  -9.827 < 2e-16
## PhoneServiceYes   -0.57854    0.14494  -3.991 6.57e-05
## InternetServiceFiber optic  0.84120    0.10752   7.824 5.13e-15
## InternetServiceNo  -0.72024    0.14942  -4.820 1.43e-06
## OnlineSecurityYes  -0.49034    0.09413  -5.209 1.90e-07
## ContractOne year   -0.67067    0.11803  -5.682 1.33e-08
## ContractTwo year   -1.39744    0.19416  -7.198 6.13e-13
## PaperlessBillingYes  0.36526    0.08213   4.447 8.71e-06
## PaymentMethodCredit card (automatic) -0.08679    0.12595  -0.689 0.49081
## PaymentMethodElectronic check  0.33125    0.10477   3.162 0.00157

```

```
## PaymentMethodMailed check      -0.17305    0.12792   -1.353    0.17612
## TotalCharges                    0.87972    0.15954    5.514 3.50e-08
##
## (Intercept)                    ***
## SeniorCitizen1                  ***
## tenure                          ***
## PhoneServiceYes                  ***
## InternetServiceFiber optic      ***
## InternetServiceNo                ***
## OnlineSecurityYes                ***
## ContractOne year                 ***
## ContractTwo year                 ***
## PaperlessBillingYes              ***
## PaymentMethodCredit card (automatic)
## PaymentMethodElectronic check    **
## PaymentMethodMailed check
## TotalCharges                     ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6553.1  on 5648  degrees of freedom
## Residual deviance: 4715.9  on 5635  degrees of freedom
## AIC: 4743.9
##
## Number of Fisher Scoring iterations: 6
```

```
vif(df_clean.bestlogit)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## SeniorCitizen    1.080908  1      1.039667
## tenure           14.617958  1      3.823344
## PhoneService      1.401480  1      1.183841
## InternetService    2.402199  2      1.244951
## OnlineSecurity     1.127523  1      1.061849
## Contract           1.521245  2      1.110580
## PaperlessBilling   1.113757  1      1.055347
## PaymentMethod      1.343729  3      1.050474
## TotalCharges       16.492869  1      4.061141
```

```
getinfo(df_clean.bestlogit,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]
```

```
## $confusion_matrix
##      predicted
## observed  0   1
##      No  773 248
##      Yes  82 280
##
```

```
## $accuracy
## [1] 0.7613883
##
## $sensitivity
## [1] 0.7734807

# Remove Total Charges due to high VIF value (>2, thus multi-collinearity effect)

df_clean.bestlogit2 <- glm(Churn~
                           SeniorCitizen +
                           tenure +
                           PhoneService +
                           InternetService +
                           OnlineSecurity +
                           Contract +
                           PaperlessBilling +
                           PaymentMethod,
                           family = binomial,
                           data = df_clean2[training==1,])
summary(df_clean.bestlogit2)

##
## Call:
## glm(formula = Churn ~ SeniorCitizen + tenure + PhoneService +
##      InternetService + OnlineSecurity + Contract + PaperlessBilling +
##      PaymentMethod, family = binomial, data = df_clean2[training ==
##      1, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8218  -0.6731  -0.3068   0.7653   3.1154
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.31487    0.15882  -8.279  < 2e-16
## SeniorCitizen1     0.37046    0.09132   4.057 4.98e-05
## tenure          -0.78160    0.05612 -13.927  < 2e-16
## PhoneServiceYes   -0.40597    0.13731  -2.957 0.003111
## InternetServiceFiber optic  1.07104    0.09861  10.862  < 2e-16
## InternetServiceNo -0.85564    0.14626  -5.850 4.91e-09
## OnlineSecurityYes -0.44134    0.09379  -4.706 2.53e-06
## ContractOne year  -0.63739    0.11659  -5.467 4.58e-08
## ContractTwo year  -1.30311    0.19059  -6.837 8.07e-12
## PaperlessBillingYes  0.37475    0.08173   4.585 4.53e-06
## PaymentMethodCredit card (automatic) -0.08768    0.12574  -0.697 0.485626
## PaymentMethodElectronic check   0.35246    0.10472   3.366 0.000763
## PaymentMethodMailed check    -0.11915    0.12660  -0.941 0.346609
##
```

```

## (Intercept) ***
## SeniorCitizen1 ***
## tenure ***
## PhoneServiceYes **
## InternetServiceFiber optic ***
## InternetServiceNo ***
## OnlineSecurityYes ***
## ContractOne year ***
## ContractTwo year ***
## PaperlessBillingYes ***
## PaymentMethodCredit card (automatic)
## PaymentMethodElectronic check ***
## PaymentMethodMailed check
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6553.1  on 5648  degrees of freedom
## Residual deviance: 4748.6  on 5636  degrees of freedom
## AIC: 4774.6
##
## Number of Fisher Scoring iterations: 6

```

```

vif(df_clean.bestlogit2)

```

```

##              GVIF Df GVIF^(1/(2*Df))
## SeniorCitizen    1.081676  1      1.040037
## tenure           1.662398  1      1.289340
## PhoneService     1.358571  1      1.165577
## InternetService  1.924459  2      1.177815
## OnlineSecurity   1.113718  1      1.055328
## Contract         1.481751  2      1.103300
## PaperlessBilling 1.108524  1      1.052865
## PaymentMethod    1.318029  3      1.047098

```

```

getinfo(df_clean.bestlogit2,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]

```

```

## $confusion_matrix
##      predicted
## observed  0  1
##      No  780 241
##      Yes  87 275
##
## $accuracy
## [1] 0.7628344
##
## $sensitivity
## [1] 0.7596685

```

```

# Decision Tree Analysis
df_clean.fulltree <- rpart(Churn ~.,
                           data = df_clean2[training==1,], method = "class",
                           control = rpart.control(cp=0))

getinfo(df_clean.fulltree,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]

## $confusion_matrix
##      predicted
## observed    0    1
##      No   811 210
##      Yes  127 235
##
## $accuracy
## [1] 0.7563268
##
## $sensitivity
## [1] 0.6491713

# Hyperparameter Tuning

# plotcp(df_clean.fulltree)
printcp(df_clean.fulltree)

##
## Classification tree:
## rpart(formula = Churn ~ ., data = df_clean2[training == 1, ],
##       method = "class", control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] Contract      Dependents      DeviceProtection gender
## [5] InternetService MonthlyCharges MultipleLines OnlineBackup
## [9] OnlineSecurity PaperlessBilling Partner      PaymentMethod
## [13] PhoneService SeniorCitizen StreamingMovies StreamingTV
## [17] TechSupport tenure      TotalCharges
##
## Root node error: 1507/5649 = 0.26677
##
## n= 5649
##
##      CP nsplit rel error  xerror    xstd
## 1 0.07011723     0  1.00000 1.00000 0.022058
## 2 0.01360319     3  0.78965 0.79761 0.020412
## 3 0.00398142     5  0.76244 0.81287 0.020553
## 4 0.00265428    10  0.73723 0.80027 0.020437
## 5 0.00248839    18  0.71267 0.79429 0.020381
## 6 0.00232250    22  0.70272 0.79429 0.020381
## 7 0.00199071    31  0.67750 0.79496 0.020387

```

```
## 8 0.00176952    32 0.67551 0.79695 0.020406
## 9 0.00165893    49 0.63504 0.79695 0.020406
## 10 0.00149303   61 0.61447 0.79628 0.020400
## 11 0.00132714   65 0.60849 0.80226 0.020455
## 12 0.00110595   75 0.59456 0.80624 0.020492
## 13 0.00099536   81 0.58792 0.80823 0.020510
## 14 0.00088476   91 0.57664 0.80956 0.020523
## 15 0.00082946  107 0.55740 0.80956 0.020523
## 16 0.00079628  113 0.55209 0.81221 0.020547
## 17 0.00066357  120 0.54612 0.82681 0.020679
## 18 0.00049768  126 0.54214 0.82681 0.020679
## 19 0.00044238  137 0.53417 0.84472 0.020838
## 20 0.00033179  140 0.53285 0.84472 0.020838
## 21 0.00026543  142 0.53218 0.85468 0.020924
## 22 0.00022119  147 0.53086 0.86662 0.021027
## 23 0.00016589  150 0.53019 0.86662 0.021027
## 24 0.00013271  159 0.52820 0.86662 0.021027
## 25 0.00000000  164 0.52754 0.86662 0.021027
```

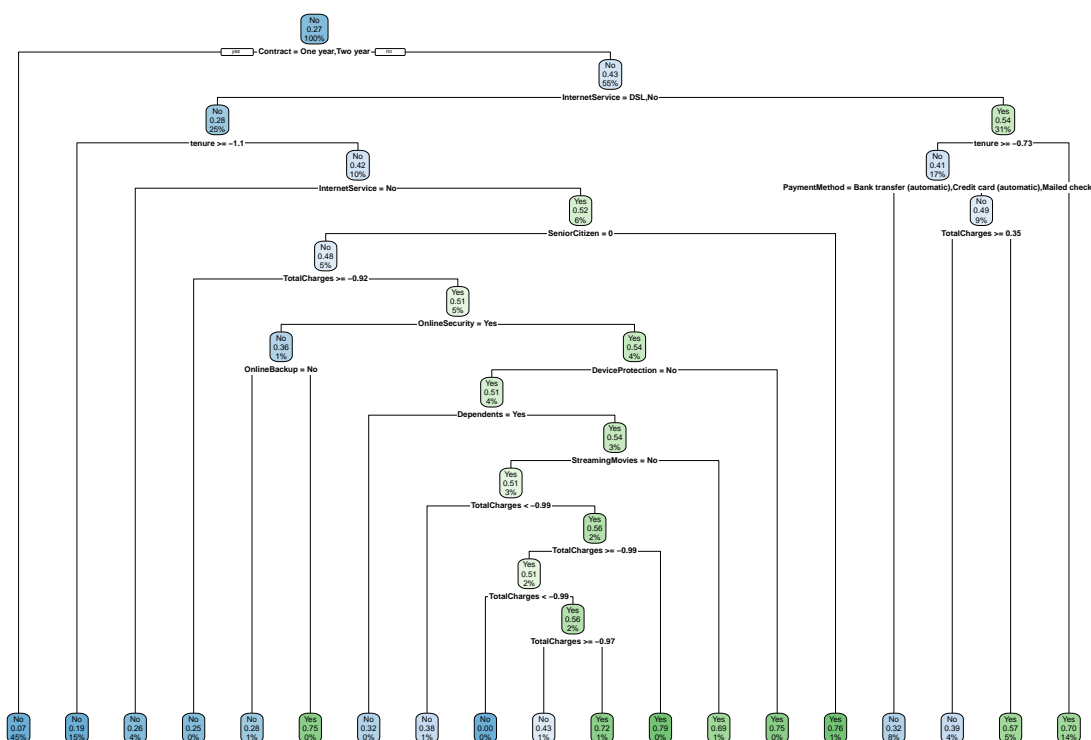
```
tmp <- df_clean.fulltree$cptable[which.min(df_clean.fulltree$cptable[, "xerror"]),]
```

```
# Prune the tree
```

```
df_clean.besttree <- prune(df_clean.fulltree, cp = tmp[1])
```

```
rpart.plot(df_clean.besttree)
```





```
getinfo(df_clean.besttree,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]
```

```
## $confusion_matrix
##      predicted
## observed  0   1
##      No  814 207
##      Yes  106 256
##
## $accuracy
## [1] 0.7736804
##
## $sensitivity
## [1] 0.7071823
```

```
# Random Forest
set.seed(1994)
df_clean.rforest <- randomForest(Churn~.,
                                data = df_clean2[training==1,],
                                ntree=500,                      # dataset
                                cutoff=c(0.5,0.5),
                                mtry=2,
                                importance=TRUE)

df_clean.rforest
```


```
##
## Call:
##  randomForest(formula = Churn ~ ., data = df_clean2[training ==      1, ], ntree = 500, cut
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 2
##
##      OOB estimate of  error rate: 20.46%
## Confusion matrix:
##      No Yes class.error
## No  3789 353  0.08522453
## Yes   803 704  0.53284672

# Confusion Matrix Test
getinfo(df_clean.rforest,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]

## $confusion_matrix
##      predicted
## observed    0    1
##      No   837 184
##      Yes  119 243
##
## $accuracy
## [1] 0.7809111
##
## $sensitivity
## [1] 0.6712707

# Hyperparameter Tuning
set.seed(1994)
rforest.tune <- tuneRF(x = df_clean2[training==1,]%>%select(-Churn),
                      y = df_clean2[training==1,]$Churn,mtryStart=2,
                      ntreeTry = 500)

## mtry = 2  OOB error = 20.38%
## Searching left ...
## mtry = 1  OOB error = 22.61%
## -0.10947 0.05
## Searching right ...
## mtry = 4  OOB error = 20.62%
## -0.01216334 0.05
```



```
# SVM
df_clean.svm <- svm(Churn~.,
                    data = df_clean2[training==1,],
                    kernel = "linear",
                    cost = 0.01,
                    probability = TRUE)

getinfo(df_clean.svm, df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]
```

```
## $confusion_matrix
##      predicted
## observed  No Yes
##      No  919 102
##      Yes  179 183
##
## $accuracy
## [1] 0.7968185
##
## $sensitivity
## [1] 0.5055249
```

```
# Hyperparameter Tuning
svm.tune <- tune(svm,
                 Churn~.,
```

```

        data = df_clean2[training==1,],
        kernel = "linear",
        ranges = list(cost = 10^(-5:0)))

print(svm.tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.199325

svm.tune$best.model

##
## Call:
## best.tune(method = svm, train.x = Churn ~ ., data = df_clean2[training ==
##   1, ], ranges = list(cost = 10^(-5:0)), kernel = "linear")
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  0.1
##       gamma: 0.03225806
##
## Number of Support Vectors: 2608

df_clean.bestsvm <- svm(Churn~.,
                        data = df_clean2[training==1,],
                        kernel = "linear",
                        cost = 0.1,
                        probaility = TRUE)

getinfo(df_clean.bestsvm,df_clean2)[c("confusion_matrix", "accuracy", "sensitivity")]

## $confusion_matrix
##      predicted
## observed  No  Yes
##      No  915 106
##      Yes 178 184
##
## $accuracy
## [1] 0.7946493

```


```
##
## $sensitivity
## [1] 0.5082873

# Performance evaluation - Learning Curves and Fitted Graphs
# AUC Curve
# First assemble the probability matrix
prob_matrix <- data.frame(
  "logit" = predict(df_clean.bestlogit2, df_clean2[training==2,], type = "response"),
  "d_tree" = predict(df_clean.besttree, df_clean2[training==2,], type="prob")[,2],
  "r_forest" = predict(df_clean.rforest, df_clean2[training==2,], type = "prob")[,2],
  "svm" = as.numeric(attr(predict(df_clean.bestsvm, df_clean2[training==2,], decision=2)))
)

# Create the ROC Variable

logit.roc <- roc(df_clean2$Churn[training==2], prob_matrix$logit)
d_tree.roc <- roc(df_clean2$Churn[training==2], prob_matrix$d_tree)
r_forest.roc <- roc(df_clean2$Churn[training==2], prob_matrix$r_forest)
svm.roc <- roc(df_clean2$Churn[training==2], prob_matrix$svm)

ggroc(list(logit=logit.roc, decision.tree=d_tree.roc, random.forest=r_forest.roc, svm=svm.roc), legend="none")
```



```
tmp4 <- c(logit.roc$auc, d_tree.roc$auc, r_forest.roc$auc, svm.roc$auc)
```

```
tmp5 <- data.frame(
  "AUC" = tmp4
)
row.names(tmp5)<- c("Logit", "Decision Tree", "Random Forest", "SVM")
```

```
tmp5
```

```
##              AUC
## Logit        0.8358193
## Decision Tree 0.8087443
## Random Forest 0.8314863
## SVM          0.8270775
```

A small discussion about cutoff point:

As we are attempting to identify customer that are going to churn, we thus need to focus on sensitivity metric compared to accuracy. As it is comparatively more expensive to acquire customer than retain customer, thus we are not as concern with false positive, but rather concerned with false negative. We would ideally like a model that is able to successful target all customer that are going to churn, and it should matter less if we have a higher number of false postive to us a telcommunication company. Thus we should have a lower threshold value than 0.5, though the actual value often require domain knowledge which we lack, thus we are going to use a more objective method to set out threshold value.

```
# Lets use logistic regression as it has the largest AUC out of all three method
```

```
# output <- matrix(0,100,3)
# x_axis <- seq(0.01,0.8,length=100)
#
# for (i in 1:100)
# {
#   output[i,]=threshold(x_axis[i])
# }
#
# plot(x_axis,output[,1], type = "l", col = "darkgreen", xlab = "Threshold Value", ylab = "Val
# lines(x_axis,output[,2],col = "red")
# lines(x_axis,output[,3], col = "blue")
# legend("bottom",col=c(2,"darkgreen",4,"darkred"),text.font =3,inset = 0.02,
#       box.lty=0,cex = 0.8,
#       lwd=c(2,2,2,2),c("Specificity", "Senitivity", "Accuracy"))
#
#
# x_axis[which(abs(output[,1]-output[,2])<0.01)]
```